


End-to-End Pipeline Dataflow: Ingests raw signals across 14 sources (4 core + 5 job boards + 5 news feeds) via protocol router, structures data through resilient chunked LLM extraction, resolves entities via 4-tier candidate matching, and syncs canonical records to Google Sheets.



1. Scale Strategy: Autonomous Collection of 500,000+ Records

Current Implementation (Verified Proof Point): The pipeline executes an asynchronous, non-blocking core (`asyncio.gather`) governed by a centralized protocol router (`src/scrapper/router.py`). Concurrency and politeness policies are decoupled from scraping logic and driven entirely by declarative configurations (`config.yaml`), guaranteeing **zero hardcoded pagination or volume limits** (verified by scaling YC to 1,427 startups, SaaSHub to 1,050 products, and Arxiv/PwC to 1,149 papers without manual intervention):

- **Granular Domain Semaphores:** Each domain operates under an isolated `asyncio.Semaphore` paired with enforced pacing (e.g., `concurrency: 5`, `delay: 1.0s` for SaaSHub and YC; `concurrency: 1`, `delay: 15.0s` for Arxiv). Shared semaphores ensure concurrent workers collectively respect domain boundaries without burst violations.
- **Dual-Engine Protocol Routing:** Fast endpoints run over `aiohttp` connection pooling, while bot-protected or dynamic JavaScript targets (e.g., MIT News encountering Akamai TLS fingerprint blocks) automatically route to an async headless `Playwright` Chromium cluster with shared connection reuse (`src/scrapper/browser_client.py`).
- **Stateful Resumption:** The crawlers query existing SQLite primary keys before issuing pagination calls, skipping known entity ranges and deduplicating at the ingestion boundary.

Scaling to 500,000+ Entities Without Human Intervention:

1. **Distributed Worker Topology (Celery / Temporal + Redis):** Replace single-process `asyncio.gather` with distributed task queues. A discovery scheduler partitions the crawl space and enqueues cursor-based pagination tasks into Redis. Stateless worker nodes across multi-region containers pull partition chunks independently.
2. **Global Coordination via Distributed Token Bucket:** In-memory Python semaphores fail across distributed nodes. Production scale enforces per-domain rate limits globally using a **Redis-backed Token Bucket (executed via atomic Lua scripts)**. Before issuing any HTTP request, a worker requests a token for `ratelimit:{domain}`. If empty, the worker yields back to the queue with a computed delay, guaranteeing strict domain-level compliance across 100+ parallel workers.
3. **Automated Discovery Breadth:** For Arxiv, ingest via monthly category partitions (`cs.AI`, `cs.LG`, `cs.CL`, `cs.CV`) using resume tokens (`resumptionToken`) over OAI-PMH harvest protocols. For software directories, parallelize breadth-first directory traversal across 120+ category slugs. For startups, federate ingestion to OpenCorporates and SEC EDGAR Form D filings.
4. **Adaptive Backpressure & Circuit Breakers:** Workers track rolling HTTP status codes per domain. If an elevated 403/429 ratio occurs, an automated circuit breaker trips, dynamically backing off crawl velocity or rotating egress IP subnets (residential proxy mesh) without human intervention.

2. Handling 413s & 429s: Context Windows & High-Concurrency Rate Limiting

Current Implementation (Verified Proof Point): The extraction subsystem (`src/extraction/`) structures unstructured payloads (Hacker News comments, RSS HTML bodies, entity arbitration pairs) under strict LLM API limits, achieving 100% extraction success with zero unhandled 429s across 4,800+ requests:

- **Semantic Payload Chunking** (`src/extraction/chunking.py`):
 - *Hard Ceiling*: Payloads are capped at 4,000 input tokens (~12,000 chars), reserving 1,000+ tokens for instructions and JSON output schemas to eliminate **HTTP 413 (Payload Too Large)**.
 - *Inverted-Pyramid News Chunking*: Front-loads essential facts from lede and context while concluding with the takeaway. The chunker preserves the title, first 3 paragraphs, and final paragraph, excising low-density middle prose with an explicit `[...content truncated for brevity...]` token.
 - *Header-Preserving Job Parsing*: HN hiring comments isolate structured metadata in line 1 (`Company | Title | Location | Remote`). The chunker preserves line 1 intact and captures the first 500 chars of the body (tech stack and compensation), discarding administrative boilerplate.
- **Resilient Retry Policy with Tenacity** (`src/extraction/llm_client.py`): Rate limits are mitigated via exponential backoff with random additive jitter: `Wait = min(16.0, initial * 2^attempt) ± jitter`, configured with `initial=1.0s`, `multiplier=2.0`, `max=16.0s`, and `jitter=±0.5s` over `max_retries=3`. Jitter breaks concurrency synchronization locks across workers, preventing "thundering herd" retry spikes on provider gateways.
- **Three-Tier Provider Failover Chain**: Groq (Qwen 3.8 27B) -> Gemini (Gemini 3.6 Flash) -> OpenRouter (DeepSeek V3). If Groq exhausts its retry budget (or trips its 1,000 output-token/min free quota), the request cascades to Gemini without throwing an unhandled exception or dropping records.

Scaling to Thousands of Concurrent Extractions:

1. **Hierarchical Extraction (Deterministic First, LLM Second)**: At 500k records, running an LLM on every record is cost-prohibitive (> \$1,500) and induces severe rate limits. The architecture mandates **Strategy A (Deterministic Fallback)**: fast C-based regex parsers and CSS selectors extract 90%+ of structured fields at microsecond latency. The LLM failover chain is only invoked for genuine nulls, ambiguous roles, or cross-lingual titles (verified in `test_fallback_parsers.py`).
2. **Asynchronous Batch API Processing**: Non-real-time ingestion (paper abstracts, startup summaries) routes to OpenAI/Anthropic/Gemini **Batch APIs** (50% cost reduction, dedicated 24-hour throughput pools immune to interactive rate limits).
3. **Local Embedding Pruning**: Entity resolution candidate pairs bypass LLM arbitration entirely when embedding cosine similarities are > 0.92 (deterministic merge) or < 0.70 (deterministic reject). Only the 70%–92% borderline band is dispatched to LLM arbitration, reducing LLM calls by 98.4%.

3. Freshness Tracking: Zero Duplicate Processing Across Distributed Nodes

Current Implementation (Verified Proof Point): The Phase II crawler enforces rolling 24-hour UTC freshness across 5 news feeds and 5 job boards via a multi-tiered validation pipeline (`src/scrapper/date_normalizer.py`):

- **Tier 1 & Tier 2 (Strict Temporal Filters)**: Normalizes ISO-8601, RFC-822, and relative strings ("2 hours ago") to UTC. Records older than 24.0 hours are dropped immediately at the ingestion boundary.
- **Tier 3 (Persistent Deduplication Registry)**: For sources lacking publication dates (e.g. un-timestamped listings), URLs are checked against a persistent SQLite `seen_urls` table storing (`url` TEXT PRIMARY KEY, `first_seen_at` TIMESTAMP). If a URL already exists, it is recognized as historical and discarded.
- **Compound Storage Uniqueness**: `startups` enforces `UNIQUE(source_url, entity_name)`; `products` enforces `UNIQUE(source_url, startup_name)`; `research_papers` enforces `UNIQUE(paper_url)` (version suffixes `v1`, `v2` stripped to deduplicate Arxiv vs PwC); `news / jobs` enforce `UNIQUE(url) / UNIQUE(job_url)`.

Scaling to Distributed Crawler Nodes:

1. **Distributed Two-Tier URL Registry (Redis Bloom Filter + PostgreSQL)**:
 - *Tier A: In-Memory Redis Scalable Bloom Filter*: Crawler nodes test URL presence in O(1) time with zero disk I/O. With false positive probability $p = 0.001$, a 500,000-URL filter consumes only ~900 KB of RAM. If the Bloom filter returns false, the URL is guaranteed unseen.
 - *Tier B: Distributed Lease Locks* (`SETNX url_hash`): Before fetching, a node acquires a 10-minute lease lock on the URL hash in Redis. If Worker B encounters the same URL while Worker A is actively crawling it, Worker B skips it immediately, preventing concurrent in-flight network duplicates.
2. **Event-Driven Partitioning via URL Hashing**: The crawl frontier is partitioned by consistent hashing: `Partition = MurmurHash3(canonical_domain) % N`. All URLs for a specific domain route to the exact same worker group, enforcing both domain rate limits and freshness checks locally without cross-node network chatter.
3. **Conditional HTTP Revalidation (ETag / If-Modified-Since)**: Periodic freshness audits for news feeds issue lightweight HTTP `HEAD` or conditional `GET` requests with cached `ETag` and `If-Modified-Since` headers. Unmodified feeds return `304 Not Modified`, consuming 0 parsing or LLM compute.

4. Storage Strategy: Primary Database, Vector Indexing, and Graph Modeling

Justification of Current Architecture (Trial Scale: <10,000 Records): For this trial's scope (1,000+ records per category), **SQLite** (`data/records.db`) was selected deliberately over a hosted client-server database: zero network overhead and microsecond C-library access; strict ACID compliance and atomic multi-table transactions; zero external infrastructure dependencies (single portable file enabling frictionless replication and offline grading); and built-in compound **UNIQUE** constraints enforcing relational data integrity.

Production Storage Architecture at 500,000+ Scale:

1. **Primary Relational Store: PostgreSQL (with Partitioning):** Multi-Version Concurrency Control (MVCC) eliminates SQLite table-level write locks under high multi-worker concurrency. Connection pooling is managed via **PgBouncer**. Tables are partitioned by month and source (`research_papers_2026_q3`, `jobs_2026_w36`), enabling parallel query scans and fast archival. **JSONB** columns store heterogeneous source payloads with GIN indexes on nested metadata.
2. **Vector Store: PostgreSQL + pgvector (HNSW Indexing):** Entity descriptions, taglines, and paper abstracts are vectorized via `text-embedding-3-small` (1536 dimensions). Replaces lexical inverted token indexing with **Hierarchical Navigable Small World (HNSW)** vector indexing (`distance = cosine`, `m = 16`, `ef_construction = 64`), enabling sub-millisecond retrieval of top-k nearest neighbors across 500,000 items to form candidate pairs for arbitration.
3. **Graph Database: Neo4j or Amazon Neptune:** Models multi-hop cross-domain relationships that are computationally intractable in SQL `JOIN` tables: `(Company)-[:MAKES]->(Product)`, `(Company)-[:POSTS]->(Job)`, `(Author)-[:AUTHORS]->(Paper)-[:CITES]->(Code)`. Transitive entity resolution (e.g. `Merge A~B and B~C => A~C`) executes via native Graph Connected Component algorithms rather than iterative SQL self-joins, preventing circular clustering loops.

Engineering Judgment: Known Tradeoffs & Production Remediations

Observed Limitation	Trial Mitigation	500k+ Production Remediation
<code>seen_urls</code> First-Run Blind Spot	Accepted for initial batch; strict 24h filter active on subsequent runs.	Pre-seed crawl registry from Internet Archive CDX API and historical sitemap archives.
<code>SaaS</code> Hub <code>startupName</code> = <code>productName</code>	Phase IV entity resolution clusters product extensions under parent brands.	Implement dual-source parent company enrichment using Clearbit/Brandfetch Company APIs.
<code>monday.com</code> Domain Suffix Tradeoff	Enables high-confidence fuzzy matching against product extensions (<code>monday CRM</code>).	Adopt dual-representation schema: normalized stem keys for matching, authoritative registry for display.

Quantitative Architecture Summary & Validation Metrics

Architecture Component	Current Pipeline Implementation	500,000+ Record Target Architecture
Ingestion & Concurrency	<code>aiohttp</code> + Playwright, per-domain Semaphore (1,427 startups, 1,050 products, 1,149 papers)	Distributed Celery/Temporal on K8s + Redis Token Bucket (Atomic Lua)
LLM Orchestration	Groq (<code>Qwen 3.8</code>) -> Gemini -> OpenRouter (100% failover success, 0 unhandled 429s)	Strategy A deterministic filters + Async Batch API
Entity Resolution	Inverted Token Index (99.8% pair reduction) + 4-Tier Matching (100% false-friend rejection)	HNSW Vector Indexing (<code>pgvector</code> , k-NN) + Cross-Encoder
Primary Storage & Sync	SQLite (<code>records.db</code> , compound UNIQUE) + <code>gsread</code> 6-tab live synchronization	Partitioned PostgreSQL (PgBouncer) + Neo4j Graph + Parquet Lakehouse